

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/343903563>

Benchmarking Semi-supervised Federated Learning

Preprint · August 2020

DOI: 10.48550/arXiv.2008.11364

CITATION

1

READS

2,600

6 authors, including:



Yaoqing Yang

Dartmouth College

77 PUBLICATIONS 3,167 CITATIONS

SEE PROFILE



Yujun Yan

Dartmouth College

18 PUBLICATIONS 1,136 CITATIONS

SEE PROFILE



Joseph E. Gonzalez

University of California, Berkeley

188 PUBLICATIONS 6,888 CITATIONS

SEE PROFILE



Michael W. Mahoney

University of California, Berkeley

359 PUBLICATIONS 23,820 CITATIONS

SEE PROFILE

Benchmarking Semi-supervised Federated Learning

Zhengming Zhang^{1*}, Zhewei Yao^{2*}, Yaoqing Yang², Yujun Yan³,
Joseph E. Gonzalez², Michael W. Mahoney^{2,4}

¹ Southeast University

² University of California, Berkeley

³ University of Michigan, Ann Arbor

⁴ International Computer Science Institute

Abstract

Federated learning promises to use the computational power of edge devices while maintaining user data privacy. Current frameworks, however, typically make the unrealistic assumption that the data stored on user devices come with ground truth labels, while the server has no data. In this work, we consider the more realistic scenario where the users have only unlabeled data and the server has a limited amount of labeled data. In this semi-supervised federated learning (SSFL) setting, the data distribution can be non-iid, in the sense of different distributions of classes at different users. We define a metric, R , to measure this non-iidness in class distributions. In this setting, we provide a thorough study on different factors that can affect the final test accuracy, including algorithm design (such as training objective), the non-iidness R , the communication period T , the number of users K , the amount of labeled data in the server N_s , and the number of users $C_k \leq K$ that communicate with the server in each communication round. We evaluate our SSFL framework on Cifar-10, SVHN, and EMNIST. Overall, we find that a simple consistency loss-based method, along with group normalization, achieves better generalization performance, even compared to previous supervised federated learning settings. Furthermore, we propose a novel grouping-based model average method to improve convergence efficiency, and we show that this can boost performance by up to 10.79% on EMNIST, compared to the non-grouping based method.

1 Introduction

Current state-of-the-art machine learning models can potentially benefit from the large amount of user data privately-held on mobile devices, as well as the computing power locally-available on these devices. In response to this, federated learning (FL), which only requires transmitting the trained (intermediate) models, has been proposed as a privacy-preserving solution to exploit the data and computing power on mobile devices [1, 2]. In a typical FL pipeline, a server maintains a model and shares it with users/devices. Each user/device updates the global shared model for multiple steps locally using only locally-held data, and then it uploads the updated model back to the server. After aggregating all the models from users, the server takes an averaging step over all the models (e.g., FedAvg [2]), and it then sends the averaged model back to users [1, 3]. This approach respects privacy in the (weak) sense that the server does not access the private user data at any point in the procedure. However, prior work in FL has made the unrealistic assumption that the data stored on the local device are fully annotated with ground-truth labels and that the server does not have access to any labeled data. In fact, the private data at the local device are more often unlabeled, since annotating data requires both time and domain knowledge [4, 5], and servers are often hosted by organizations that do have labeled data.

*equal contribution

Motivated by this, in this paper, we study a more realistic FL setting, which we call the *semi-supervised federated learning* (SSFL) setting. In SSFL, users only have access to unlabeled data, and the server only has a small amount of labeled data.¹ In this context, our main contributions are the following.

1. We introduce the SSFL framework, and we propose a principled way to study non-iid data distributions. Specifically, we introduce a probability distance measurement (total variation distance, denoted as R) to evaluate the difference in user class distributions. This is often called the *class distribution skew*. By synthesizing datasets with different R values, we are able to study SSFL under different levels of class distribution skew using quantitative evaluations.²
2. We provide an empirical evaluation of the effects of different normalization (batch normalization (BN) [11] versus group normalization (GN) [12]) and different choices of training objectives (traditional self-training loss versus consistency regularization loss (CRL) [13, 14]). Our result shows that GN combined with CRL is the best combination in terms of testing accuracy.
3. With the best solution (GN with CRL), we extensively evaluate our SSFL framework concerning: different non-iidness R ; the number of users K ; the communication period T ; the number of labeled data points in the server N_s ; and the number of users $C_k \leq K$ that communicate with the server in each communication round. In addition to establishing state-of-the-art results, we also benchmark the more realistic SSFL setting, as compared to previous work.
4. We propose a novel grouping-based model averaging method to reduce communication when the number of users is large. Particularly, we divide users into different sub-groups, and we use different FedAvg [2] models for different sub-groups. This grouping-based method outperforms FedAvg by 0.55%/0.16%/10.79% on Cifar-10/SVHN/EMNIST, respectively.
5. We compare our proposed method with other supervised FL algorithms and supervised distributed training algorithms. Even when our communication frequency is lower than the supervised distributed training algorithms (0.80%/0.29% better than EASGD/OverlapSGD [15, 9]) and/or the degree of our non-iidness is higher than other supervised FL methods, our approach still achieves superb results (12.83%/9.53% better than FedAvg/DataSharing [7, 2]).

Overall, by formulating the SSFL setting, by studying different algorithm/model designs, and by evaluating the best design under different factors (e.g., the class distribution skew R), we provide a valid baseline on SSFL.³ We have open-sourced our SSFL framework.⁴

2 Related Work

Federated Learning. Federated learning (FL) [1, 2, 16, 17, 7, 18, 19, 20, 5, 21, 22] is a decentralized computing framework that enables multiple users to learn a shared model while potentially protecting privacy of users (although recent work [23] shows this may not be the case). Federated Averaging (FedAvg) [2], which is the most popular FL algorithm, shows good performance when the data distribution across users is iid. However, in the non-iid case, the performance can significantly degrade. In fact, dealing with non-iid distributions is deemed by many to be one of the most critical challenges in FL [7, 18, 16]. In [7], a data-sharing method is proposed to improve the final accuracy. However, sharing massive data among all users requires both large storage space as well as stable connections between users and the server. Importantly, all of these methods require the data stored by the local users to come with ground-truth labels (in order to perform model updates locally). The FL problem in the semi-supervised setting, when users do not have labels, however, is “relatively ignored” and has “little prior arts,” as mentioned in a recent survey paper [22].

In addition to the challenge of the non-iidness of the data distribution and the need for local ground truth labels, communication efficiency is another critical problem in FL [24, 19, 25, 26, 27]. One way to

¹As we completed our manuscript, we became aware of a very recent paper introducing similar ideas [6].

²Synthesizing datasets with skewed class distributions has been adopted to study non-iidness in [7, 8, 9]. See Section 3.1 in [10] for a thorough discussion on different types of non-iid distributions beyond class distribution skew.

³Our approach of using CRL combined with GN and grouping-based model averaging can also be applied to the (more common) supervised FL setting. A detailed exploration of this is outside the scope of this paper, and it is likely of less interest than the (more realistic) setting we consider.

⁴<https://github.com/jhcknazzm/SSFL-Benchmarking-Semi-supervised-Federated-Learning>

relieve the communication burden of FL is to increase the period (the number of local gradient descent iterations) between consecutive communication stages. However, when this communication period increases, the diversity between different models increases, and the fusion of these models by the server may lead to accuracy degradation. To handle this problem, [25] proposes FedProx, which adds a proximal term in the user local loss function to restrict the update distance between the local model and the global model. Other work considers gradient compression and model compression to reduce the communication cost [19, 21, 27]. For example, [19] proposes atomic sparsification of stochastic gradients, which leads to significantly faster distributed training.

Regarding the motivation of SSFL, a recent survey paper [22] raises the practical concern that users may not have ground-truth labels. Regarding the problem formulation, [5, 20] are most relevant to our work. However, unlike [5], the data at the users in SSFL are fully unlabeled, and one cannot use a supervised learning technique to train the users' data; and unlike [20], we consider the non-iid case with unlabeled user data, and we study the effect of the changing number of communicating users, which is also more realistic.

Semi-supervised Learning. Semi-supervised learning (SSL) is a classical problem when only a small fraction of data is labeled [28, 4, 29, 30, 31, 32, 14, 33]. In [33], a self-training method is introduced, which improves the state-of-the-art accuracy on ImageNet [34], even compared to supervised learning [35, 36]. In [14], a consistency-based loss is proposed to improve the performance of SSL further.

In this work, we study SSL in the federated setting. The main challenges, compared to standard SSL, are two-fold: (i) the non-iid data distribution; and (ii) the communication constraints. In particular, for the non-iid data distribution: all of the labeled data are only stored in the server, as opposed to classical SSL setting where labeled data are shared with all computational nodes; and the data stored in different local users have an unbalanced class distribution.

3 A Framework of Semi-supervised Federated Learning

3.1 Our semi-supervised federated learning framework

In this section, we provide a thorough study on different factors that can affect the test accuracy of SSFL. In our setup, there exists a cloud server and K users/devices. We denote the labeled dataset at the server by $D_s = \{(x_i, y_i)\}_{i=1}^{N_s}$, and the unlabeled dataset stored at the k -th user by $D_k = \{x_i\}_{i=1}^{N_k}$, for $k \in \{1, \dots, K\}$. Here, N_s (N_k) is the number of labeled (unlabeled) samples available at the server (k -th user). Note that no data are exchanged between the server and the users. That is to say, the server can only train on dataset D_s , and the k -th user can only train the local model using the local dataset D_k . Our basic SSFL setup is illustrated in Fig. 1. In this work, we consider image classification as a representative SSFL task.

Similar to traditional FL [2, 1], in our SSFL framework, there are two key components, i.e., the local training process and how to combine those locally-trained models. For the first component (the local training process), our SSFL framework has two main difference from the traditional FL: on the users' side, they do not have the labeled data. Therefore, new training techniques need to be introduced (see § 3.3); on the server's side, instead of only averaging the models from users as traditional FL does, our server also needs to train its own model. For the model combination process, since the server also has its solely trained model in SSFL, the server needs to compute an averaged model using all the received models (including its own).

We discuss different training approaches and variants of the averaging methods (referred to as algorithm design) in § 3.3. Besides those specific algorithm design (as well as model design, like normalization method as shown in next section), in this work, we extensively explore the different general factors (independent to algorithm/model design) affecting the performance of SSFL. In particular, the following factors are considered,

1. Non-iidness R : the non-iid metric of data distributions; see Definition 1.
2. Communication period T : during two consecutive communications, the number of training steps for local users and the server.
3. User number K : the total number of users.

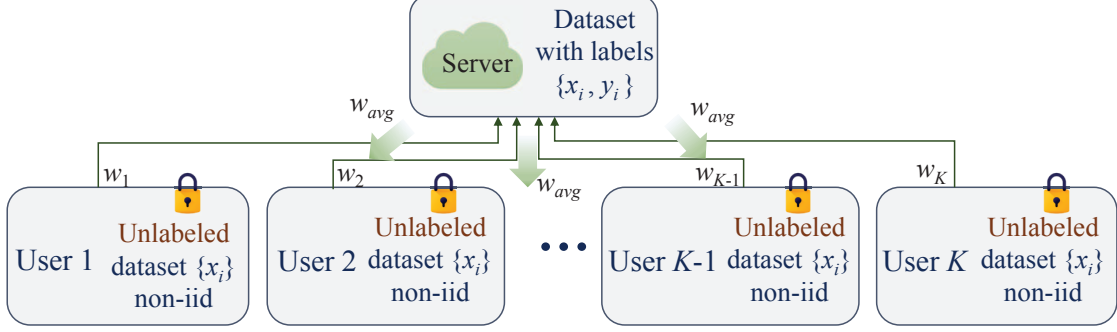


Figure 1: Semi-supervised federated learning (SSFL). Only the server has access to labeled data, i.e., the data stored in local users are unlabeled. Furthermore, the data distributions across different users are non-iid.

4. Server data number N_s : the number of labeled data in the server (since the entire dataset is fixed, the remaining data will be distributed among the client devices).
5. The group of participating clients C_k : at each communication, the group of clients who send their models to the server.

3.2 A metric on the level of non-iidness

We study the non-iidness in the sense of class distribution skew [7, 8, 9], e.g., a single user can have more data for one (couple) class(es) than others. To formally define such skew in SSFL, we use the average total variation distance described in the following definition. In the definition, the empirical class distribution of the data D_k at the k -th user is denoted by $P_k \in \mathbb{R}^d$, where d is the number of classes.⁵

Definition 1 (Metric R for non-iid level). *The non-iid metric R to measure the class distribution skew is defined as:*

$$R = \frac{1}{K(K-1)} \sum_{1 \leq k < m \leq K} \|P_k - P_m\|_1, \quad (1)$$

where $\|\cdot\|_1$ is the L_1 norm.

Observe that Definition 1 can be expressed as

$$R = \frac{1}{K(K-1)/2} \sum_{1 \leq k < m \leq K} \|P_k - P_m\|_1/2,$$

where $\|P_k - P_m\|_1/2$ is the (normalized) total variation distance, which takes value in $[0, 1]$, and where $K(K-1)/2$ is the number of user pairs, i.e., it is the mean total variation distance averaged over pairs of users. In particular, $0 \leq R \leq 1$ [37]. When data are distributed in such a way that each user has a uniform empirical class distribution $P_k = [1/d, \dots, 1/d]$, for all k , we have $R = 0$; and in another extreme, when $K = d$ and each user only has access to one class, we have $R = 1$.

Remark 1. *The metric R in Definition 1 does not explicitly consider the effect of different data sizes N_k 's at different users. We focus on the case when N_k 's are equal to each other, while slight difference may arise when the overall number of samples is not divisible by the number of users.*

To study the impact of non-iid levels, we need to synthesize datasets with different R values. The specific data distribute procedures to achieve R are relegated to § A.2.

⁵Here, $\sum_{j=1}^d P_k[j] = 1$, for all $1 \leq k \leq K$.

3.3 Algorithm and model design

Here, we discuss specific choices on the algorithm and model designs that we consider.

Training Objective. We denote the loss function used by the server (resp., k -th user) as L_s (resp., L_k). Inspired by the traditional self-training method [38, 33], we can define L_s and L_k as follows:

$$L_s = \frac{1}{N_s} \sum_{(x_i, y_i) \in D_s} l(y_i, f_s(\alpha(x_i); w_s)), \quad (2)$$

$$L_k = \frac{1}{N_k} \sum_{x_i \in D_k} \mathbf{1}(\max(\bar{y}_i) \geq \tau) l(\arg \max(\bar{y}_i), f_k(\alpha(x_i); w_k)), \quad (3)$$

where w_s (w_k) are the weights of server model f_s (k -th user model f_k), $l(\cdot, \cdot)$ is the cross-entropy loss, $\alpha(\cdot)$ is the data augmentation function (for traditional self-training, it is flip-and-shift augmentation), \bar{y}_i is the prediction of the model f_k on the augmented sample $\alpha(x_i)$, $\mathbf{1}(\cdot)$ is the indicator function, and τ is the threshold hyperparameter, which helps the model to decide which samples have high confidence to be trained. A simple and popular method to optimize Eq. 2 and Eq. 3 is SGD with momentum over a mini-batch.

In [14], the authors use two different types of data augmentations (DA): the standard flip-and-shift augmentation $\alpha(\cdot)$ (referred to as *weak DA*); and the RandAugment [39] $A(\cdot)$ (referred to as *strong DA*). Here, the latter RandAugment uses two different augmentation methods (i.e., shift and crop) out of twelve possible augmentation methods (e.g., rotate, shift, solarize, etc.) for one image. We refer the interested reader to [39] for a detailed explanation. The key idea behind using two DAs (i.e., weak DA and strong DA) is that the predictions of the same image with two data augmentations should be consistent with each other. Recall that, on the user side, the data have no label. Therefore, using this approach, we can use the pseudo-labels generated from weak DA samples to train strong DA samples. This can boost the testing performance due to the large diversity of samples generated by strong DA. We refer Eq. 4 as the consistency regularization loss (CRL),

$$L_k = \frac{1}{N_k} \sum_{x_i \in D_k} \mathbf{1}(\max(\bar{y}_i) \geq \tau) l(\arg \max(\bar{y}_i), f_k(A(x_i); w_k)), \quad (4)$$

where $\bar{y}_i = f_k(\alpha(x_i); w_k)$. This CRL has proved effective in the SSL setting [14], but its impact in SSFL is unexplored. Also, recall that the SSL setting is different from our SSFL setting: (i) both labeled data and unlabeled data are used for each training iteration (this can be viewed as the sum of Eq. 2 and 4); (ii) there is no non-iidness; and (iii) there is usually no constraint on the user-server communication.

Model Averaging. A naive way to perform model averaging is using FedAvg [2],

$$w_{avg} \stackrel{FedAvg}{=} \left(w_s + \sum_{k \in C_k} w_k \right) / (|C_k| + 1). \quad (5)$$

After this, the server will broadcast this w_{avg} to all the users for next round of training. As show in § 4.3, if we use naive FedAvg, given in Eq. 5, the large model diversity across different users significantly slows down the training process. To overcome this, we propose a novel grouping-based model averaging, which divides C_k communication users into S groups and then performs the average group-wise. Specifically, after collecting all C_k communication users, the server divides them into S groups $\{G_i\}_{i=1}^S$, and updates the w_{avg} according to:

$$\begin{cases} w_{avg,i} = \left(w_s + \sum_{w_k \in G_i} w_k \right) / (|G_i| + 1), & i \in \{1, 2, \dots, S\} \\ w_{avg} = \sum_{i=1}^S w_{avg,i} / S. \end{cases} \quad (6)$$

After this, $w_{avg,i}$ is broadcast to the user group G_i , and w_{avg} is the weights used for server training.

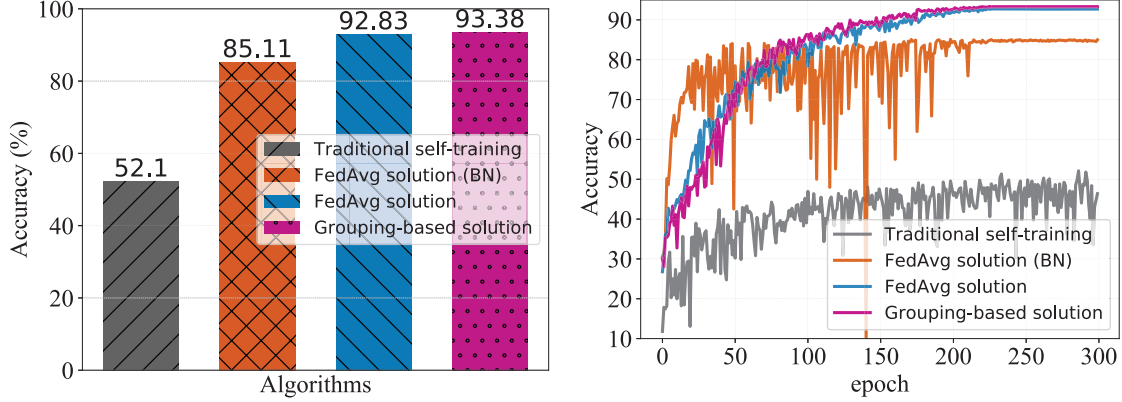


Figure 2: (Left) Accuracy of different methods on Cifar-10. (Right) The corresponding convergence behaviors.

Normalization. Since the data distributions across users are non-iid in SSFL, normalization methods (e.g., BN) using statistical properties of data distributions, will have statistical discrepancies between different users. This can lead the average model to obtain *worse* generalization performance. Recent works [40, 41] found out that in supervised FL with non-iid data distributions, the performance of group normalization (GN) is usually much better than that of batch normalization (BN). Here, we also study these two different normalization methods (BN versus GN) in our SSFL setting.

Experimental comparison of different factors. We use ResNet-18 [35] on Cifar-10 to study the effectiveness of different choices (loss, averaging methods, and normalization scheme) mentioned above. Here, for the rest of the factors mentioned in § 3.1, we set $T = 16$, $K = 10$, $C_k = 10$, $R = 0.4$, and $N_s = 1000$. The threshold τ used in Eq. 2 and Eq. 4 is chosen to be 0.95, as in [14]. The number of groups S in Eq. 6 is set as 2. In particular, we abbreviate:

1. “Traditional self-training”: we use Eq. 3; BN; and FedAvg, Eq. 5.
2. “FedAvg solution (BN)”: we use Eq. 4; BN; and FedAvg, Eq. 5.
3. “FedAvg solution”: we use Eq. 4; GN; and FedAvg, Eq. 5.
4. “Grouping-based solution”: we use Eq. 4; GN; and grouping-based model averaging, Eq. 6.

In Fig. 2, we compare the results of different approaches. As one can see, without any modification of the model/loss/averaging method, the accuracy of traditional self-training is only 52.10%. Replacing the user loss from Eq. 3 to Eq. 4, the performance boosts to 85.11%. Further correcting the normalization method from BN to GN helps the accuracy increase to 92.83%. The best performance is achieved by our grouping-based model averaging, which has 0.55% better accuracy than the FedAvg solution. The convergence behaviors of different methods are also shown in the right of Fig. 2.

Later on, without specification, the results reported in all of our experiments use the FedAvg solution setting. (A detailed comparison between FedAvg and group-based solution is discussed in § 4.3.)

4 Empirical Evaluation

We consider three datasets, Cifar-10, SVHN, and EMNIST in our empirical evaluation. We use ResNet-18 [35] as the training model on both Cifar-10 and SVHN datasets; and we use the same CNN model as [40] on EMNIST. For EMNIST, we use the “balanced” training dataset. That is, we truncated the training dataset to have 2400 data points per class and drop the rest. For the test dataset, we use the full test set without truncating. See Appendix A for more details.




Figure 3: (Left) Comparison between different non-iid levels (R) on Cifar-10, SVHN and EMNIST. (Middle) Accuracy versus communication period T . (Right) Accuracy versus labeled data points in the server (N_s).

4.1 Impact of non-iidness R , communication period T , and labeled data number N_s

We conduct experiments on Cifar-10, SVHN and EMNIST to illustrate the effect of the non-iid level R . Particularly, for Cifar-10/SVHN/EMNIST, we set user number $K = 10/10/47$, the amount of labeled data $N_s = 1000/1000/4700$, the number of communicate users $C_k = 10/10/10$, and the communication period $T = 16/16/16$, respectively. The results are shown in the left of Fig. 3. When $R = 0$ ($R = 1$), the user will have a uniform empirical class distribution (single class data). As can be seen, the accuracy decreases as the non-iid level R increases (from 93.38%/94.15%/84.15% to 79.39%/86.34%/82.56% on Cifar-10/SVHN/EMNIST). This is in accord with our intuition that iid data distribution typically leads to the best result.

We also illustrate the effect of the communication period T on Cifar-10, SVHN and EMNIST. Particularly, we set $R = 0.4/0.4/0.4$, $K = 10/10/47$, $N_s = 1000/1000/4700$ and $C_k = 10/10/10$ for Cifar-10, SVHN and EMNIST, respectively. The middle of Fig. 3 presents our results. Increasing T (i.e., communicating less frequently) leads to a worse generalization performance. This can be explained since the local model can be overfitted when T is large. In addition, the converges behaviors on Cifar-10 at different T can be found in Fig. B.1.

To investigate the impact of the amount of labeled data in the server, we design our experiment on Cifar-10, SVHN and EMNIST. Particularly, we use $T = 16/16/16$, $R = 0.4/0.4/0.4$, $K = 10/10/47$ and $C_k = 10/10/10$ in all experiments. Our result is shown in the right of Fig. 3. Note that increasing the amount of labeled data in the server can significantly improve final generalization performance. For example, with 4000 labeled data, the accuracy on Cifar-10/SVHN/EMNIST is 1.69%/2.99%/2.7% higher as compared to 1000 labeled data. This is intuitive since the increase of the amount of labeled data can make the model trained by the server more accurate, which helps the users to obtain more accurate pseudo-labels and to improve performance. In the extreme case, if the server has the entire labeled training dataset, the situation degrades to a centralized supervised learning setting.

4.2 Impact of the number of communicated users C_k

In the real scenario, the number of connected users varies at different time due, e.g., to the internet quality or privacy concerns. To simulate this drop-and-reconnected case, for every communication, we here assume that only $C_k \leq K$ users out of the entire K users are connecting to the server. Namely, for every communication, only C_k users will upload their local trained model into the server. This setting also has another advantage, i.e., that the communication volume remains the same as we increase the user number K , while holding C_k fixed.

We set $T = 16/16/16$, $R = 0.4/0.4/0.4$, and $N_s = 1000/1000/4700$ for the three datasets (Cifar-10, SVHN, and EMNIST). The results of our FedAvg method on Cifar-10 and SVHN are shown in Tab. 1, and

Table 1: Accuracy versus amount of communicated user C_k on Cifar-10 and SVHN. Here, “*” here means we train SVHN for $E = 120$ epochs instead of $E = 40$ epochs for normal SVHN training.

	$K = 10, C_k = 10$	$K = 20, C_k = 20$	$K = 30, C_k = 30$
Cifar-10	92.83%	92.13%	91.87%
SVHN	94.36%	94.83%	74.71% (93.94%*)

	$K = 10, C_k = 10$	$K = 20, C_k = 10$	$K = 30, C_k = 10$
Cifar-10	92.83%	93.06%	92.78%
SVHN	94.36%	94.78%	93.37%

Table 2: The results of FedAvg setting versus grouping-based setting on Cifar-10, SVHN and EMNIST.

Dataset	FedAvg	Grouping-based
Cifar-10	92.83%	93.38%
SVHN	94.83%	94.99%
EMNIST	70.84%	81.63%

the result of EMNIST is presented in Tab. C.1.

On the top of Tab. 1, we set $C_k = K$. As can be seen, increasing the number of users has a marginal effect ($<1\%$) on the final performance, from $K = 10$ to $K = 30$. One notable thing here is that with $K = 30$, if we use the standard training epochs (40 epochs) on SVHN, it only has 74.71% accuracy, which is 19.65% lower than $K = 10$. With a larger number of users, the diversity of the models across different users increases significantly. This will add noise to the federated averaging and will make the convergence slow (discussed below). The training curve (as shown in Fig. D.1) verifies our conjecture. Therefore, we increase the training epochs from 40 to 120 for $K = 30$ on SVHN, and the final accuracy is 93.94%.

The bottom of Tab. 1 shows the result with fixed $C_k = 10$ and various K (from 10 to 30). Counterintuitively, the results have consistently better performance, compared to $C_k = K$. Particularly, the $K = 30, C_k = 10$ case outperforms $C_k = 30$ by 0.91%/18.66% on Cifar-10/SVHN, respectively.

As mentioned above ($K = C_k = 30$ case on SVHN), this phenomenon can be explained by too much diversity of models across users. Here, using $C_k < K$ helps reduce the diversity. In order to explore our conjecture, inspired by [42], we use weight diversity to analyze the diversity of models across users. The weight diversity⁶ is defined as:

$$\Delta_{gd}^t(w) = \sum_{i=1}^K \|w_i^t - w_{avg}^t\|_2^2 / \sum_{i=1}^K (w_i^t - w_{avg}^t)\|_2^2, \quad (7)$$

where w_*^t represents the weights at the t -th iteration for the server/user. We conduct our weight diversity experiments on the classifier layer of ResNet-18 on Cifar-10, and the results in the setting of $t \bmod T = 0$ are shown in Fig. 4. Note that when $C_k = K$, the convergence speed of $K = 30$ is much slower than $K = 20$ (red curves in the top of Fig. 4). However, with fixed $C_k = 10$, the effect of increasing K from 20 to 30 is negligible (purple curves in the top of Fig. 4). The bottom of Fig. 4 show the weight diversity under different settings. It is obvious that the diversity of $C_k = 10$ is much less, compared to $C_k = K$ for both $K = 20$ and $K = 30$. From a finer-grained view, when $C_k = K$, the weight diversity (purple curves) does not vanish until epoch 20. However, for $C_k < K$, the weight diversity (red curves) quickly converges to a flat region. This is consistent with the training curve shown on the top of Fig. 4.

4.3 Impact of model averaging methods

To verify the effectiveness of our grouping-based model averaging algorithm, we perform experiments on all three datasets. In particular, we set $K = 10/20/47$, $R = 0.4/0.4/0.4$, $C_k = 10/20/47$, $T = 16/16/16$, and $N_s = 1000/1000/4700$ for Cifar-10/SVHN/EMNIST, respectively. For the grouping-based algorithm, we use group number $S = 2/2/5$ on Cifar-10/SVHN/EMNIST, respectively.

Our results are shown in Tab. 2. It can be seen that the accuracy of the grouping-based setting is 0.55%/0.16%/10.79% higher than FedAvg on three datasets. Note that our grouping-based structure can reduce the effect of the $C_k = K$ problem (i.e., weight diversity issue), as discussed in § 4.2 on EMNIST.

⁶Here, we can think of this weight diversity as the same as gradient diversity [42] since multi-step updates in a local user can be considered as the gradient accumulation.




Figure 4: (Top) The convergence of $K = 20/30$ on Cifar-10. (Bottom) The weight diversity of $K = 20/30$ on Cifar-10.

See the EMNIST results in Tab. 2 and Tab. C.1. Besides, the convergence behavior of the grouping-based method can be found in Appendix E.

4.4 Comparison with supervised results

Table 3: Comparison with supervised FL. Our proposed framework achieves better results even when the non-iid level is high and users have no labeled data. Here, “*” is calculated according to the setting in DataSharing [7].

Scenario	Supervised FedAvg [2] (users have labels)	DataSharing [7] (users have labels)	Our FedAvg solution (users have no label)
Cifar-10, Non-iid	78.52% ($R = 0.29$)	81.82% ($R = 0.29^*$)	91.35% ($R = 0.4$)

We compare our FedAvg method with other FL algorithms in Tab. 3. Due to the lack of SSL setting in the current literature, we choose two supervised FL methods for comparison, Supervised FedAvg [2] and DataSharing [7]. We set $K = 10$, $C_k = 10$ and $T = 32$, and we use ResNet-18 [35] to be the model for training. The non-iid setting of DataSharing [7] corresponds to the scenario where we set $R = 0.29$. For our FedAvg solution, we set $N_s = 1000$ and $R = 0.4$. From Tab. 3 we see that the performance of the our FedAvg method ($R = 0.4$) on Cifar-10 is still better than supervised FedAvg ($R = 0.29$) and DataSharing methods ($R = 0.29$).

Table 4: Comparison with EASGD [15] and OverlapSGD [9] on Cifar-10.

Scenario	EASGD [15] (users have labels)	OverlapSGD [9] (users have labels)	Our FedAvg solution (users have no label)	Our Grouping-based solution (users have no label)
Non-iid, $T = 2$	91.12%	91.63%	93.83%	94.22%
Non-iid, $T = 8$	88.88%	91.45%	92.52%	93.58%
Non-iid, $T = 32$	—	—	91.28%	91.92%

We also compare our solutions (both FedAvg and grouping-based model averaging) with EASGD [15] and OverlapSGD [9] which are communication efficient algorithms under supervised settings. We set $K = 16$, $R = 0.4$, $C_k = 16$ and $N_s = 1000$ on Cifar-10 for our methods, and the results are shown in Tab. 4. As one can see, our result is much better than both EASGD and OverlapSGD. Particularly, even with $T = 32$, our grouping-based solution has 0.80%/0.29% better performance, as compared to EASGD/OverlapSGD in the setting of $T = 2$, respectively. Note that both EASGD and OverlapSGD are supervised algorithms, which means they have all the data labeled.

5 Conclusions

We have proposed the first semi-supervised federated learning (SSFL) framework, where the server has a small amount of labeled data, and the users have only unlabeled data. We defined a metric to measure quantitatively the non-iid level; and we studied our framework under different factors, including the non-iidness R , the communication period T , the number of users K , the number of labeled data N_s , and the number of users C_k communicating to the server with each communication round, in order to benchmark our SSFL framework. Our extensive empirical evaluation conducted on Cifar-10, SVHN, and EMNIST demonstrated that the simple consistency loss-based method, along with group normalization, could already achieve higher generalization performance than previous supervised FL settings. In addition, with our novel grouping-based model averaging, we can further improve generalization performance. It is worth emphasizing that our method has a certain generality, and it can be easily extended to other FL scenarios, such as traditional FL [2] and label-centralized and distributed FL [4]. Another challenging issue that is worth mentioning is the case when there is significant mismatch between the user data distributions and the distribution at the server, in which case the label supervision from the server may conflict with the information provided by users. We envision that techniques from unsupervised domain adaptation could be introduced to solve this problem [43].

Acknowledgments

We would like to thank Jianyu Wang and Daniel Rothchild for their valuable feedback. We would like to acknowledge DARPA, NSF, and ONR for providing partial support of this work.

Broader Impact

Our work opens up new horizons that could lead to influential future research directions for SSFL. Researchers who work on FL or distributed learning topics may benefit from this research. With the support of this benchmark, SSFL can help realize the broader application of technology, especially in areas such as medical, financial, and legal services. These areas require a lot of time and domain knowledge to provide labels. On the other hand, the implications of SSFL are also likely to be wide: (i) mobile users can obtain intelligent image recognition services without providing any labeled data, on the other hand, this may also cause mobile users to consume a lot of power; and (ii) a model trained on biased data may continue the inherent bias in the SSFL setting. We will conduct further research work to understand the biases and limitations of the datasets used in SSFL.

References

- [1] J. Konečný, H. B. McMahan, F. X. Yu, P. Richtárik, A. T. Suresh, and D. Bacon, “Federated learning: Strategies for improving communication efficiency,” *arXiv preprint arXiv:1610.05492*, 2016.
- [2] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, *et al.*, “Communication-efficient learning of deep networks from decentralized data,” *arXiv preprint arXiv:1602.05629*, 2016.
- [3] J. Konečný, B. McMahan, and D. Ramage, “Federated optimization: Distributed optimization beyond the datacenter,” *arXiv preprint arXiv:1511.03575*, 2015.
- [4] X. Zhu and A. B. Goldberg, “Introduction to semi-supervised learning,” *Synthesis lectures on artificial intelligence and machine learning*, vol. 3, no. 1, pp. 1–130, 2009.
- [5] A. Albaseer, B. S. Ciftler, M. Abdallah, and A. Al-Fuqaha, “Exploiting unlabeled data in smart cities using federated learning,” *arXiv preprint arXiv:2001.04030*, 2020.
- [6] W. Jeong, J. Yoon, E. Yang, and S. J. Hwang, “Federated semi-supervised learning with inter-client consistency,” *arXiv preprint arXiv:2006.12097*, 2020.
- [7] Y. Zhao, M. Li, L. Lai, N. Suda, D. Civin, and V. Chandra, “Federated learning with non-iid data,” *arXiv preprint arXiv:1806.00582*, 2018.
- [8] T.-M. H. Hsu, H. Qi, and M. Brown, “Measuring the effects of non-identical data distribution for federated visual classification,” *arXiv preprint arXiv:1909.06335*, 2019.
- [9] J. Wang, H. Liang, and G. Joshi, “Overlap local-SGD: An algorithmic approach to hide communication delays in distributed SGD,” in *ICASSP 2020-2020 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pp. 8871–8875, 2020.
- [10] P. Kairouz, H. B. McMahan, B. Avent, A. Bellet, M. Bennis, A. N. Bhagoji, K. Bonawitz, Z. Charles, G. Cormode, R. Cummings, *et al.*, “Advances and open problems in federated learning,” *arXiv preprint arXiv:1912.04977*, 2019.
- [11] S. Ioffe and C. Szegedy, “Batch normalization: Accelerating deep network training by reducing internal covariate shift,” *arXiv preprint arXiv:1502.03167*, 2015.
- [12] Y. Wu and K. He, “Group normalization,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, pp. 3–19, 2018.
- [13] M. Sajjadi, M. Javanmardi, and T. Tasdizen, “Regularization with stochastic transformations and perturbations for deep semi-supervised learning,” in *Advances in neural information processing systems*, pp. 1163–1171, 2016.
- [14] K. Sohn, D. Berthelot, C.-L. Li, Z. Zhang, N. Carlini, E. D. Cubuk, A. Kurakin, H. Zhang, and C. Raffel, “Fixmatch: Simplifying semi-supervised learning with consistency and confidence,” *arXiv preprint arXiv:2001.07685*, 2020.
- [15] S. Zhang, A. E. Choromanska, and Y. LeCun, “Deep learning with elastic averaging sgd,” in *Advances in neural information processing systems*, pp. 685–693, 2015.
- [16] X. Peng, Z. Huang, Y. Zhu, and K. Saenko, “Federated adversarial domain adaptation,” *arXiv preprint arXiv:1911.02054*, 2019.
- [17] S. Caldas, J. Konečný, H. B. McMahan, and A. Talwalkar, “Expanding the reach of federated learning by reducing client resource requirements,” *arXiv preprint arXiv:1812.07210*, 2018.

- [18] E. Jeong, S. Oh, H. Kim, J. Park, M. Bennis, and S.-L. Kim, “Communication-efficient on-device machine learning: Federated distillation and augmentation under non-iid private data,” *arXiv preprint arXiv:1811.11479*, 2018.
- [19] H. Wang, S. Sievert, S. Liu, Z. Charles, D. Papailiopoulos, and S. Wright, “Atomo: Communication-efficient learning via atomic sparsification,” in *Advances in Neural Information Processing Systems*, pp. 9850–9861, 2018.
- [20] N. Guha, A. Talwalkar, and V. Smith, “One-shot federated learning,” *arXiv preprint arXiv:1902.11175*, 2019.
- [21] J. Xu, W. Du, R. Cheng, W. He, and Y. Jin, “Ternary compression for communication-efficient federated learning,” *arXiv preprint arXiv:2003.03564*, 2020.
- [22] Y. Jin, X. Wei, Y. Liu, and Q. Yang, “Towards utilizing unlabeled data in federated learning: A survey and prospective,” 2020.
- [23] C. Xie, K. Huang, P.-Y. Chen, and B. Li, “Dba: Distributed backdoor attacks against federated learning,” in *International Conference on Learning Representations*, 2019.
- [24] K. Bonawitz, V. Ivanov, B. Kreuter, A. Marcedone, H. B. McMahan, S. Patel, D. Ramage, A. Segal, and K. Seth, “Practical secure aggregation for privacy-preserving machine learning,” in *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, pp. 1175–1191, 2017.
- [25] A. K. Sahu, T. Li, M. Sanjabi, M. Zaheer, A. Talwalkar, and V. Smith, “Federated optimization for heterogeneous networks,” *arXiv preprint arXiv:1812.06127*, vol. 1, no. 2, p. 3, 2018.
- [26] X. Yao, C. Huang, and L. Sun, “Two-stream federated learning: Reduce the communication costs,” in *2018 IEEE Visual Communications and Image Processing (VCIP)*, pp. 1–4, IEEE, 2018.
- [27] H. B. McMahan, D. M. Bacon, J. Konecny, and X. Yu, “Communication efficient federated learning,” Nov. 7 2019. US Patent App. 16/335,695.
- [28] Z.-H. Zhou and M. Li, “Tri-training: Exploiting unlabeled data using three classifiers,” *IEEE Transactions on knowledge and Data Engineering*, vol. 17, no. 11, pp. 1529–1541, 2005.
- [29] A. Rasmus, M. Berglund, M. Honkala, H. Valpola, and T. Raiko, “Semi-supervised learning with ladder networks,” in *Advances in neural information processing systems*, pp. 3546–3554, 2015.
- [30] A. Tarvainen and H. Valpola, “Mean teachers are better role models: Weight-averaged consistency targets improve semi-supervised deep learning results,” in *Advances in neural information processing systems*, pp. 1195–1204, 2017.
- [31] D. Berthelot, N. Carlini, I. Goodfellow, N. Papernot, A. Oliver, and C. A. Raffel, “Mixmatch: A holistic approach to semi-supervised learning,” in *Advances in Neural Information Processing Systems*, pp. 5050–5060, 2019.
- [32] D. Berthelot, N. Carlini, E. D. Cubuk, A. Kurakin, K. Sohn, H. Zhang, and C. Raffel, “Remixmatch: Semi-supervised learning with distribution alignment and augmentation anchoring,” *arXiv preprint arXiv:1911.09785*, 2019.
- [33] Q. Xie, E. Hovy, M.-T. Luong, and Q. V. Le, “Self-training with noisy student improves imagenet classification,” *arXiv preprint arXiv:1911.04252*, 2019.
- [34] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei, “Imagenet: A large-scale hierarchical image database,” in *2009 IEEE conference on computer vision and pattern recognition*, pp. 248–255, Ieee, 2009.

- [35] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778, 2016.
- [36] M. Tan and Q. V. Le, “Efficientnet: Rethinking model scaling for convolutional neural networks,” *arXiv preprint arXiv:1905.11946*, 2019.
- [37] M. Basseville, “Distance measures for signal processing and pattern recognition,” *Signal processing*, vol. 18, no. 4, pp. 349–369, 1989.
- [38] E. Arazo, D. Ortego, P. Albert, N. E. O’Connor, and K. McGuinness, “Pseudo-labeling and confirmation bias in deep semi-supervised learning,” *arXiv preprint arXiv:1908.02983*, 2019.
- [39] E. D. Cubuk, B. Zoph, J. Shlens, and Q. V. Le, “RandAugment: Practical data augmentation with no separate search,” *arXiv preprint arXiv:1909.13719*, 2019.
- [40] S. Reddi, Z. Charles, M. Zaheer, Z. Garrett, K. Rush, J. Konečný, S. Kumar, and H. B. McMahan, “Adaptive federated optimization,” *arXiv preprint arXiv:2003.00295*, 2020.
- [41] K. Hsieh, A. Phanishayee, O. Mutlu, and P. B. Gibbons, “The non-iid data quagmire of decentralized machine learning,” *arXiv preprint arXiv:1910.00189*, 2019.
- [42] D. Yin, A. Pananjady, M. Lam, D. Papailiopoulos, K. Ramchandran, and P. Bartlett, “Gradient diversity: a key ingredient for scalable distributed learning,” *arXiv preprint arXiv:1706.05699*, 2017.
- [43] M. Long, H. Zhu, J. Wang, and M. I. Jordan, “Unsupervised domain adaptation with residual transfer networks,” in *Advances in neural information processing systems*, pp. 136–144, 2016.
- [44] A. Krizhevsky, G. Hinton, *et al.*, “Learning multiple layers of features from tiny images,” 2009.
- [45] Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu, and A. Y. Ng, “Reading digits in natural images with unsupervised feature learning,” 2011.
- [46] G. Cohen, S. Afshar, J. Tapson, and A. Van Schaik, “EMNIST: Extending mnist to handwritten letters,” in *2017 International Joint Conference on Neural Networks (IJCNN)*, pp. 2921–2926, 2017.
- [47] I. Loshchilov and F. Hutter, “SGDR: Stochastic gradient descent with warm restarts,” *arXiv preprint arXiv:1608.03983*, 2016.
- [48] P. Goyal, P. Dollár, R. Girshick, P. Noordhuis, L. Wesolowski, A. Kyrola, A. Tulloch, Y. Jia, and K. He, “Accurate, large minibatch SGD: Training imagenet in 1 hour,” *arXiv preprint arXiv:1706.02677*, 2017.

A Additional Details on Our Empirical Evaluation.

A.1 Datasets

Cifar-10 [44] consists of images with 3 channels, each of size 32×32 pixels. Each pixel is represented by an unsigned int8. This dataset consists of 60000 color images from 10 classes, with 6000 images in each class. There are 50000 training images and 10000 test images. SVHN [45] is obtained from images of house numbers in Google Street View images. It has 99289 digits from 10 classes. There are 73257 digits for training and 26032 digits for testing. All digits have been resized to a fixed resolution of 32×32 pixels. EMNIST [46] is a set of handwritten digits which have been resized to a fixed resolution of 28×28 pixels. It is an unbalanced dataset that has 814255 digits from 62 classes, including A-Z, a-z, and 0-9. However, since the uppercase and lowercase of some handwritten letters are difficult to distinguish, for these letters, the uppercase and lowercase classes are combined into a new class. There are 15 merged letters in total, including [C, I, J, K, L, M, O, P, S, U, V, W, X, Y, Z]. Thus, there are 47 classes left. To make sure every user has almost the same amount of data, we truncated the training dataset to have 2400 data points per class and drop the

rest. The training dataset generated by the above pre-processing seeks to reduce the errors occurring from case confusion by merging all the uppercase and lowercase classes to form a class [46]. In particular, there are 112800 digits for training, and we use the full test dataset (18800 digits) for testing.

A.2 Data assignment procedures

For a dataset with d classes, to achieve a fixed R (see Definition 1), we follow the procedures below to distribute the data.

1. **(Server data)** We assign N_s labeled training samples from each of the d classes to the server. Recall that N_s is the total number of samples at the server. Thus, the class distribution at the server is uniform.⁷
2. **(User data)** Denote by n_j the number of samples left in class j after distributing the data to the server. The empirical distribution across different classes for the remaining data is $Q = [q_1, \dots, q_d]$ such that $\sum_{j=1}^d q_j = 1$, where $q_i = \frac{n_i}{\sum_{j=1}^d n_j}$. For a specific data assignment to the users, we use the *main class* of a user to refer to the class with the maximum number of samples at the user. Denote by m_j the number of users whose main class is j , and denote by $\mathcal{U}_j = \{\text{User}_1^j, \text{User}_2^j, \dots, \text{User}_{m_j}^j\}$ the set of such users. The m_j 's should be chosen to satisfy the constraint $\sum_{j=1}^d m_j = K$ in which K is the total number of users and d is the number of classes.

We use $\text{User}_1^j \in \mathcal{U}_j$ as an example to illustrate the data assignment (the same for other users):

- (a) We first assign $n_j R / m_j$ unlabeled training samples for the main class j .
- (b) After that, for each class $i \in \{1, \dots, d\}$, we assign $(1 - R)n_i q_j / m_j$ unlabeled samples, in which i can be equal to the main class j .

According to the above data assignment procedures, for a user belongs to \mathcal{U}_j (e.g., User_1^j), the number of unlabeled data for each class i is:

$$N_{j,i} = \begin{cases} n_j R / m_j + n_i q_j (1 - R) / m_j, & \text{for } i = j \\ n_i q_j (1 - R) / m_j, & \text{for } i \neq j. \end{cases} \quad (8)$$

The total number of unlabeled data of this user is $\sum_{i=1}^d N_{j,i} = n_j / m_j$. The empirical distribution $P_j = [q_{j,1}, \dots, q_{j,d}]$ across different classes can be calculated,

$$q_{j,i} = \begin{cases} R + q_j (1 - R), & \text{for } i = j \\ q_i (1 - R), & \text{for } i \neq j. \end{cases} \quad (9)$$

One can see that, for the two users whose *main* classes are j and k ($j \neq k$), we have $\|P_j - P_k\|_1 = 2R$. In this case, based on Definition 1, the total distance is R .

A.3 Optimizer

For all our computations, the optimizer we use is SGD with momentum. For the learning rate schedule, we use the cosine learning rate decay [47], shown in Eq. 10 below, which is a commonly used schedule in SSL [14]:

$$\gamma_t = \gamma \times \max \left\{ \cos \left(\pi \times c \times \frac{t - eM/B}{EM/B - eM/B} \right), \varepsilon \right\}, \quad (10)$$

where γ is the base learning rate, c is the periodic coefficient, E is the number of training epochs, B is the batch size, t is the current iteration, M is the number of training samples used in one epoch, e is the number

⁷Note that this step requires each class to have more than N_s/d samples, which is satisfied in all of our experiments.

Table A.1: *Optimizer hyperparameters used on different datasets.*

Dataset	E	M	γ	e	ε	weight decay	momentum	c	B
Cifar-10	300	65536	0.146	5	1e-4	1e-4	0.9	2.3	64
SVHN	40	65536	0.146	5	1e-4	1e-4	0.9	2.3	64
EMNIST	100	65536	0.03	0	1e-4	1e-4	0.9	0.4375	64




Figure B.1: *Test accuracy curves of our FedAvg method of different communication periods on Cifar-10.*

of epochs for warmup [48], and ε is a small constant. The hyperparameters used for different datasets can be found in Tab. A.1.

We believe that if one further tunes those hyperparameters, the performance of our FedAvg and grouping-based methods will be further improved. For example, for SVHN dataset, in the setting of $T = 16$, $R = 0.4$, $N_s = 1000$ and $K = C_k = 10$, when $E = 120$, our FedAvg method can get 96.17% accuracy, which is 2.71% higher than $E = 40$ as reported in Tab. 1.

B Convergence speed of different communication period T on Cifar-10


The accuracy curves of different communication period T on Cifar-10 are shown in Fig. B.1, which is discussed in § 4.1. The experimental settings are presented in § 4.1. From Fig. B.1, we can see that when T is small, the FedAvg method converges quickly.

C Impact of the number of communication users C_k on EMNIST dataset

This section shows the result of the impact of number of communicated users C_k on EMNIST, which is discussed in § 4.2. We set $K = 47$, $T = 16$, $R = 0.4$ and $N_s = 4700$. In this setting, the performances of our FedAvg method on EMNIST with different C_k are shown in Tab. C.1. As can be seen, a similar conclusion as we shown in § 4.2 holds. That is, when K is large, a small C_k can improve the performance.

Table C.1: Accuracy versus the amount of communicated user C_k on EMNIST dataset

	$K = 47, C_k = 10$	$K = 47, C_k = 30$	$K = 47, C_k = 47$
EMNIST	83.76%	79.05%	70.84%

**Figure D.1:** Accuracy curves of our FedAvg method for different K and E on SVHN.

D The effect of training epochs on EMNIST

In this section, we show the effect of training epochs on EMNIST when $K = C_k = 30$, which is used in Tab. 1 in § 4.2. As shown in Tab. 1, the accuracy of SVHN with $K = C_k = 30$ and $E = 40$ is much lower than $K = C_k = 10$ and $E = 40$. Increasing the number of training epochs to $E = 120$ can significantly improve the performance of $K = C_k = 30$. The comparison is shown in Fig. D.1. In the meantime, as we mentioned in Appendix A.3, increasing training epochs can also benefit $K = C_k = 10$ (the accuracy curve is also shown in Fig. D.1).

E Convergence behaviors of FedAvg and grouping-based averaging methods

In this section, we probe experiment to explore the reason why grouping-based averaging is better than FedAvg when the communication number C_k is large (Tab. 2). Similar to § 4.3, we analyze the weight diversity of the models trained by FedAvg and our grouping-based averaging method. As can be seen from Fig. E.1, when the the number of users that communicate with the server is large, the grouping-based averaging method can accelerate the convergence compared to the FedAvg method. From Fig. E.1, we also see that if we use naive FedAvg, the large model diversity across different users significantly slows down the training process. However, the grouping-based structure can effectively alleviate the weight diversity issue and improve the convergence speed.

F Impact of the ratio $\eta = C_k/K$

In real-life FL scenarios, the number of connected users can vary during training. Here, we define the ratio of connected users as $\eta = C_k/K$, and we study the impact of η on Cifar-10 and SVHN datasets.

We set $T = 16/16$, $R = 0.4/0.4$, and $N_s = 1000/1000$ for Cifar-10/SVHN. We use ResNet-18 [35] as the model for training, and we train 300/40 epochs on Cifar-10/SVHN. The results of our FedAvg method on

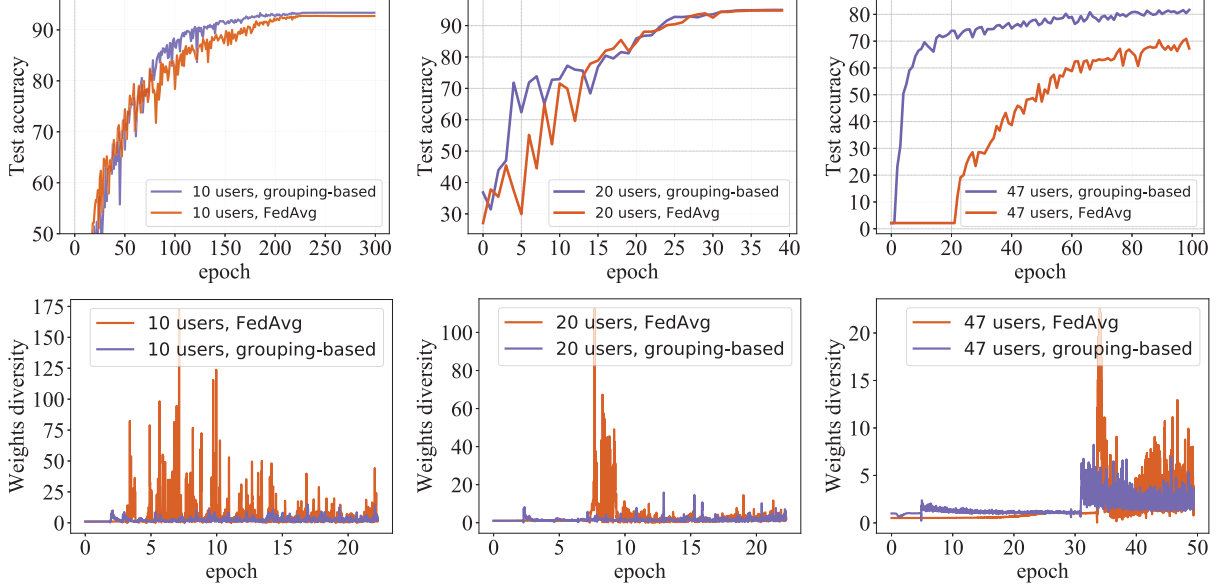


Figure E.1: (Top) The convergence behaviors of FedAvg method and grouping-based method on (left) Cifar-10, (middle) SVHN and (right) EMNIST. (Bottom) The weight diversities of FedAvg method and grouping-based method on (left) Cifar-10, (middle) SVHN and (right) EMNIST.

Table F.1: Accuracy versus different ratios of communication users on Cifar-10 and SVHN

	$\eta = 1/10$ ($K = 30, C_k = 3$)	$\eta = 1/5$ ($K = 30, C_k = 6$)	$\eta = 1/3$ ($K = 30, C_k = 10$)	$\eta = 1$ ($K = 30, C_k = 30$)
Cifar-10	80.98%	88.66%	92.78%	91.87%
SVHN	92.15%	93.12%	93.37%	74.71%

Cifar-10 and SVHN are shown in Tab. F.1. It can be found that as η increases from $\frac{1}{10}$ to $\frac{1}{3}$, the performance of our FedAvg method gradually improves. This result is intuitive since a smaller η means fewer models participate in the model averaging, only a small ratio of the models can be utilized. It should also be noted that, as shown in § 4.2 when η increases to a large value (e.g. $\eta = 1$), the diversity of models across users will be large, and the performance will decrease. Therefore, properly increasing η can improve the performance but increasing η to much can also decrease the performance. However, more work needs to be done to explore what are the most effective ratios for different datasets.

G Grouping-based averaging in fully supervised FL

In this section, we study if the grouping-based average can be extended to supervised FL (SFL). We add experiments on EMNIST under fully SFL with three different settings: $K = 47/20/10$, $T = 16/16/16$, and group number $S = 5/2/2$. The result in Table G.1 shows that the performance of the grouping method is slightly better than that of FedAvg. The performance gain of grouping-based averaging for SFL is much less than SSFL. We use the weight diversity to analyze why this is the case. The result is shown in Fig. G.1. As can be seen, it is clear that the difference of using grouping-based averaging or not is larger for SSFL than SFL.

Table G.1: The accuracy comparison of FedAvg and grouping-based average for supervised FL on EMNSIT.

User number	FedAvg	Grouping-based
$K = 47$	84.71%	84.97%
$K = 20$	86.14%	86.27%
$K = 10$	86.19%	86.29%




Figure G.1: The weight diversities of FedAvg method and grouping-based method on EMNIST when $K = 47$ for (left) SFL, and (right) SSFL. It can be clearly seen that the different between these two methods is more significant for SSFL than SFL.